

Assignment 1: COL818

SHIVAM VERMA (2020CS50442)

Implement the universality of consensus algorithm in C++

1. Implement the lock-free and wait-free variants.
2. Write your own consensus object using the compare and swap (exchange) instruction (available in both x86 and ARM). Use the `__asm` directive to implement the consensus object.
3. Implement a concurrent stack and queue using the consensus object that you created.
4. Show the correctness of your algorithm.

The submission is for machines with ARM architecture.

Files in directory

1. `LF-universalConsensus.cpp`: Lock Free variant of universal consensus object.
2. `WF-universalConsensus.cpp`: Wait Free variant of universal consensus object.
3. `Makefile`: To compile the files
4. `checker.py`: To test the correctness of execution log. To see whether property of stack and queue are maintained?
5. `execution_log.txt`: Final execution log is printed in this file, after all threads have completed their job.
6. `run.sh`: Shell script to run the task
7. `Readme`: You are currently reading this!
8. `test-consensus.cpp`: Created to test the consensus object.

How to Run?

1. In Makefile either use either SRC = LF-universalConsensus.cpp or SRC = WF-universalConsensus.cpp
2. Set number of threads in the run.sh file.
3. chmod +x run.sh
4. Run ./run.sh
5. Verification result will be printed in terminal and logs can be seen in execution_log.txt file.
6. Additional CHECK flag can be set to 1 in main file, to check which all threads are competing for a particular sequence number concurrently.

REPORT

The LFU`universal` (and WF`universal` respectively) are the main objects in the code. It contains a linked list of the nodes and a pointer to the very first node of this linked list.

The `apply` function of this object takes an `invocation` and `tid` as argument and tries to add a new node to the end of this linkedlist with this invocation using the consensus game.

Each `node` object contains a user called invocation function, sequence number and the pointer to the next node.

Consensus Function

The consensus function is implemented using a CAS protocol. We have used the `__asm` directive to implement it. The `ldaxr` and `stlxx` instructions are used.

```
template <typename T>
static inline int CAS(T *ptr, T *oldValPtr, T *newValPtr) {
    int ret;
    int res;
    T oldVal = *oldValPtr;
    T newVal = *newValPtr;
    __asm__ __volatile__ ("1:\n"
        "ldaxr %w0, [%2]\n"
        "mov %w1, %w0\n"
        "cmp %w0, %w3\n"
        "b.ne 2f\n"
        "mov %w0, %w4\n"
        "stlxx %w1, %w0, [%2]\n"
        "cbnz %w1, 1b\n"
        "2:\n"
```

```

: "=&r"(ret), "=&r"(res), "+r"(ptr)
: "r"(oldVal), "r"(newVal)
: "cc", "memory");

    return ret;
}

```

The above code is a C++ template function implementing a Compare-And-Swap (CAS) operation, typically used in concurrent programming for synchronization.

1. Input Parameters:

- **ptr**: A pointer to the memory location where the CAS operation will be performed.
- **oldValPtr**: A pointer to the old value which is expected to be found at the memory location **ptr**.
- **newValPtr**: A pointer to the new value that will replace the old value if the CAS operation is successful.

2. Local Variables:

- **ret** and **res**: These are integer variables used for return values and intermediate results in the assembly code.
- **oldVal** and **newVal**: These variables store the values pointed to by **oldValPtr** and **newValPtr**, respectively.

3. Inline Assembly Code:

- The `__asm__ __volatile__` statement introduces inline assembly code, which directly embeds assembly instructions within the C++ code.
- The assembly code implements the CAS operation. It typically loads the current value at the memory location pointed to by **ptr**, compares it with **oldVal**, and if they match, updates the memory location with **newVal**.
- It uses specific ARM assembly instructions (like `ldaxr`, `mov`, `cmp`, `stlxxr`, `cbnz`) for atomic load, store, and conditional branching, which are common in ARM architecture.

4. Output and Input Constraints:

- The `=r` constraints indicate that the corresponding variables (**ret**, **res**, **ptr**) are both inputs and outputs and will be stored in registers.
- The `r` constraints indicate that the variables **oldVal** and **newVal** will be stored in registers but are only inputs.

5. Clobber List:

- The clobber list ("cc", "memory") informs the compiler that the assembly code might modify condition flags (cc) and memory (memory), so it should not make any assumptions about their values before and after the assembly block.

6. Return Value:

- The function returns `ret`, which indicates the success of the CAS operation. Typically, it returns 0 if the operation was successful and non-zero otherwise.

STACK and QUEUE

Finally using these universal consensus object, I have implemented a concurrent `stack` and `queue` for demonstration.

Both the objects support 2 functions:

1. PUSH: push value in stack (enqueue value in queue)
2. POP: pop value from stack (deque value from the queue)

For testing, I have invoked multiple threads trying to randomly call either PUSH or POP function on these objects with equal probability.

Verification: To verify the correctness, after each thread has finished its job, the main thread makes a copy of the object and simulate the whole linkedlist of invocations on this local object. The logs are captured for each invocation and then this log is examined through a python script whether it follows `stack` (or `queue` respectively) behaviour.